

Question 1

Training a model using labeled data and using this model to predict the labels for new data is known as _____.

- a. Clustering
- b. Density estimation
- c. Unsupervised learning
- d. Supervised learning

Question 2

What are the two most common supervised tasks? (Check all that apply)

- a. Classification
- b. Reinforcement learning
- c. Clustering
- d. Regression

Question 3

Training a model using categorically labelled data to predict labels for new data is known _____.

- a. Clustering
- b. Regression
- c. Feature extraction
- d. Classification

Question 4

Training a model using labelled data where the labels are continuous quantities to predict labels for new data is known as _____.

- a. Regression
- b. Clustering
- c. Feature extraction
- d. Classification

Question 5

Modeling the features of an unlabeled dataset to find hidden structure is known as _____.

- a. Regression
- b. Unsupervised learning
- c. Supervised learning
- d. Classification

Question 6

You're running a company, and you want to develop learning algorithms to address each of two problems.

- **Problem 1:** You have a large inventory of identical items. You want to predict how many of these items will sell over the next 3 months.
- **Problem 2:** You'd like software to examine individual customer accounts, and for each account decide if it has been hacked/compromised.

Should you treat these as *classification* or as *regression* problems?

- a. Treat problem 1 as a regression problem, problem 2 as a classification problem.
- b. Treat both as regression problems.
- c. Treat both as classification problems.
- d. Treat problem 1 as a classification problem, problem 2 as a regression problem.

Question 7

Of the following examples, which would you address using an unsupervised learning algorithm?
(Check all that apply)

- a. Given a database of customer data, automatically discover market segments and group customers into different market segments.
- b. Given a dataset of patients diagnosed as either having diabetes or not, learn to classify new patients as having diabetes or not.
- c. Given email labeled as spam/not spam, learn a spam filter.
- d. Given a set of news articles found on the web, group them into set of articles about the same story.


Question 8

In a simple regression model, if you increase the input value by 1 then you expect the output to change by:

- a. The value of the *slope* parameter
- b. The value of the *intercept* parameter
- c. Also 1
- d. Impossible to tell

Question 9

Consider the following data set, and the regression line fitted on this data:




Which bold/labeled point, if removed, will have the largest effect on the fitted regression line (dashed)?

- a. b
- b. d**
- c. c
- d. a

Question 10

Using the data for classes 0, 1, and 2 plotted below, what class would a KNeighborsClassifier classify the new point as for $k = 1$ and $k = 3$?



- a. $k = 1$: Class 1; $k = 3$: Class 0
- b. $k = 1$: Class 1; $k = 3$: Class 2**
- c. $k = 1$: Class 0; $k = 3$: Class 1
- d. $k = 1$: Class 0; $k = 3$: Class 2
- e. $k = 1$: Class 2; $k = 3$: Class 1