Deep Learning

Part II

Deep = Many Hidden Layers

8 layers

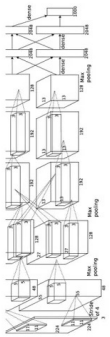

16.4%

AlexNet (2012)

19 layers

7.3%

VGG (2014)

22 layers

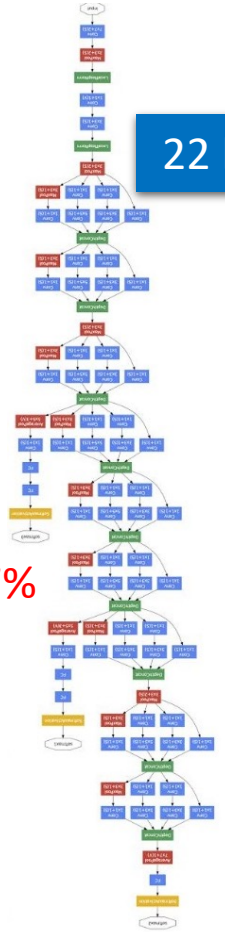

6.7%

GoogleNet (2014)

Deep = Many Hidden Layers, Continued

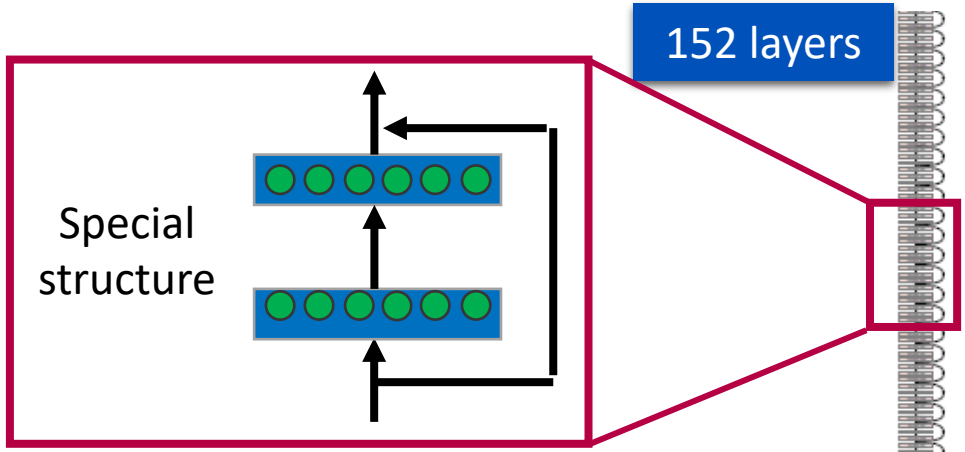

Ref:
https://www.youtube.com/watch?v=dxB6299gpvl

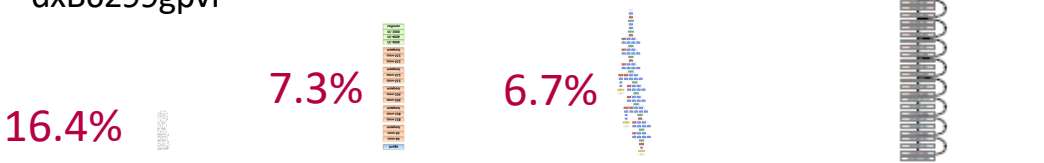

AlexNet (2012)

VGG (2014)

GoogleNet (2014)

Residual Net (2015)

Matrix Operation

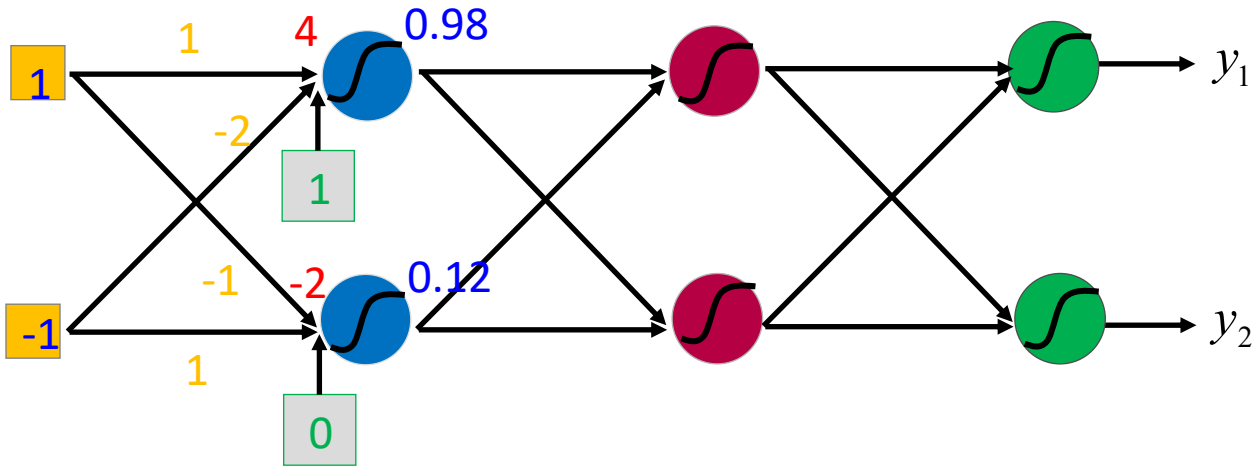

$$\sigma\left(\underbrace{\begin{bmatrix} 1 & -2 \\ -1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix}}_{\begin{bmatrix} 4 \\ -2 \end{bmatrix}} \right) = \begin{bmatrix} 0.98 \\ 0.12 \end{bmatrix}$$

Neural Network: Part II

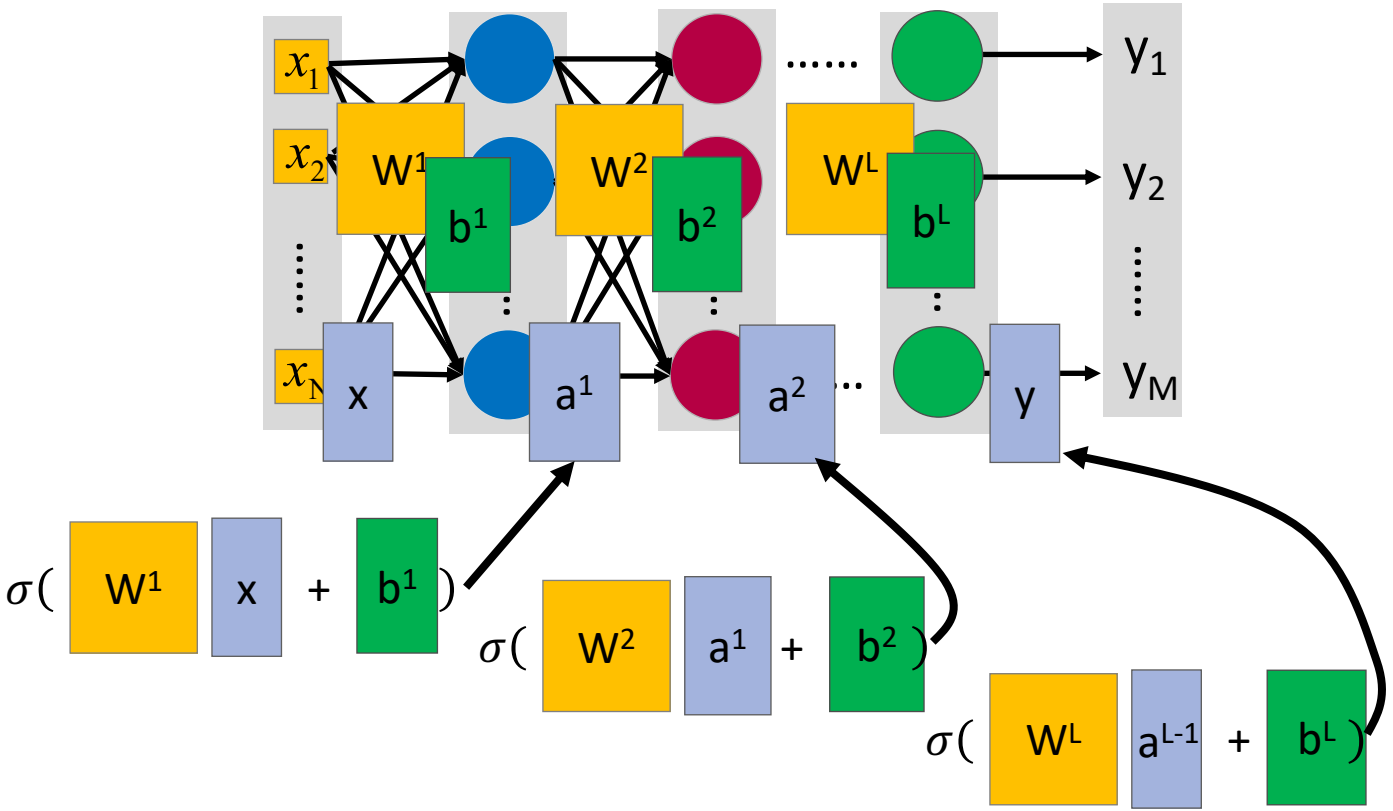

Neural Network: Part III

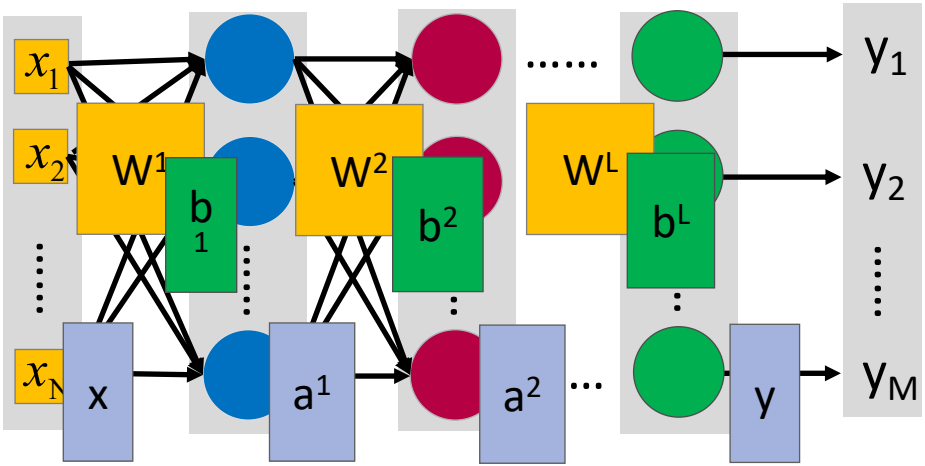

$$y = f(x)$$

Using parallel computing techniques to speed up matrix operation

$$= \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

Output Layer as Multi-Class Classifier

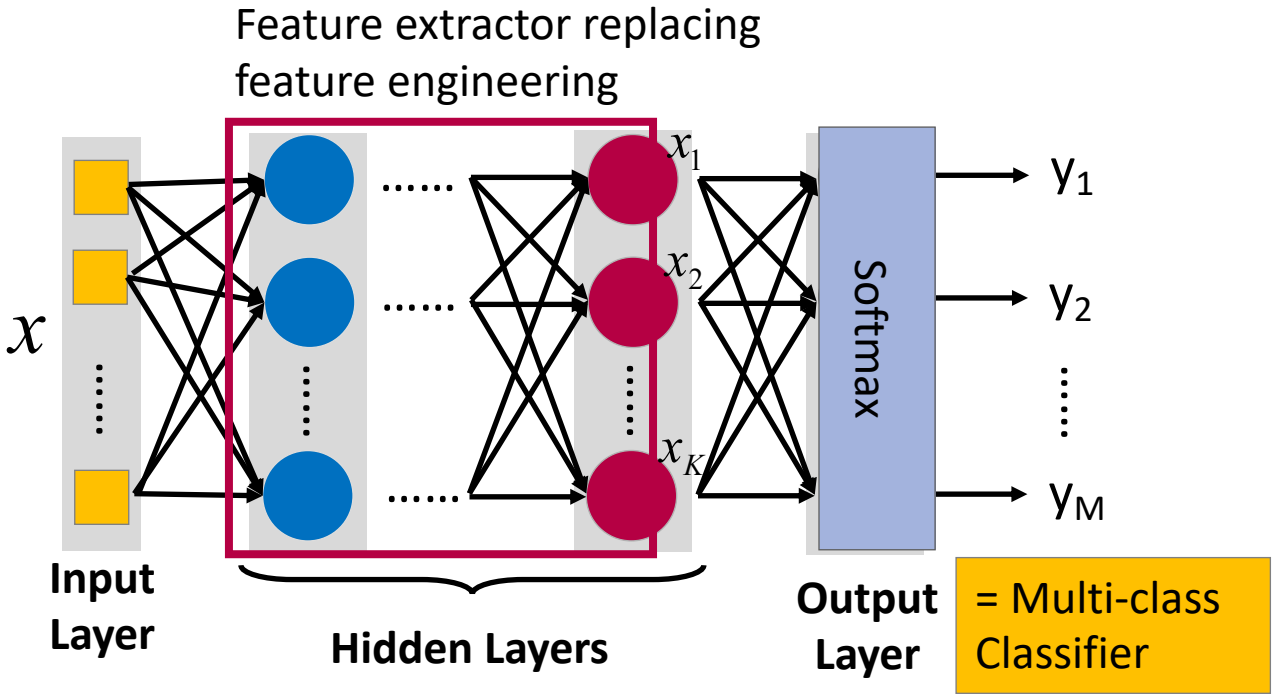

Example Application: Part I

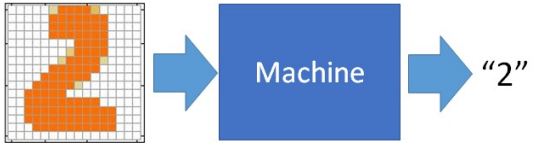

Input

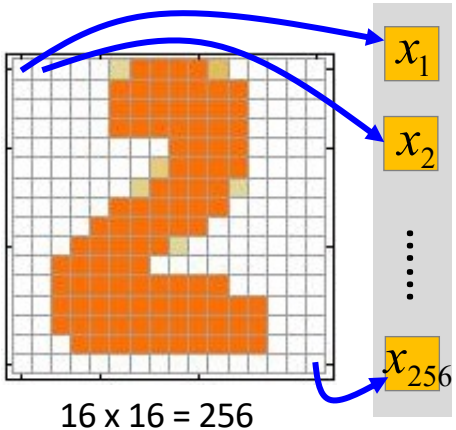

16 x 16 = 256

Ink → 1

No ink → 0

Output

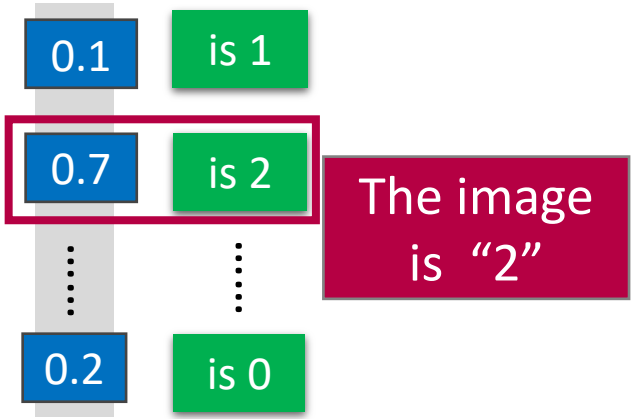

Each dimension represents the confidence of a digit.

Example Application: Part II

Handwriting Digit Recognition

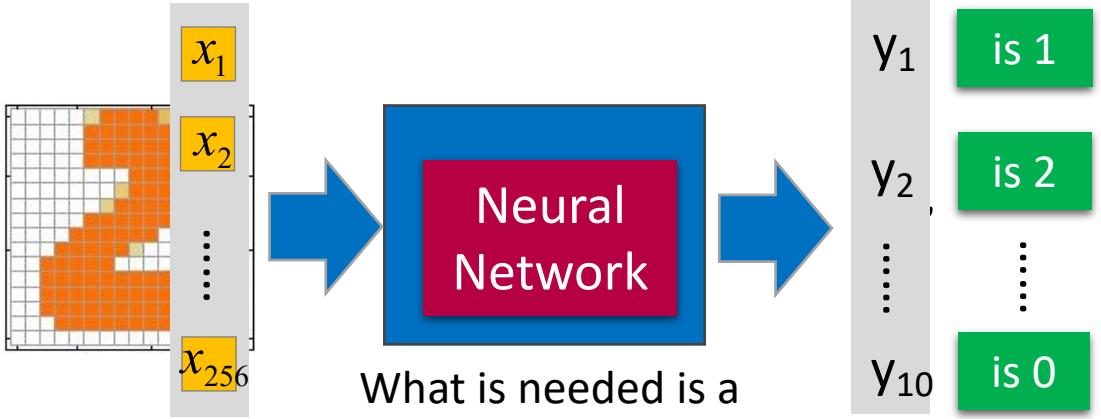

Example Application: Part III

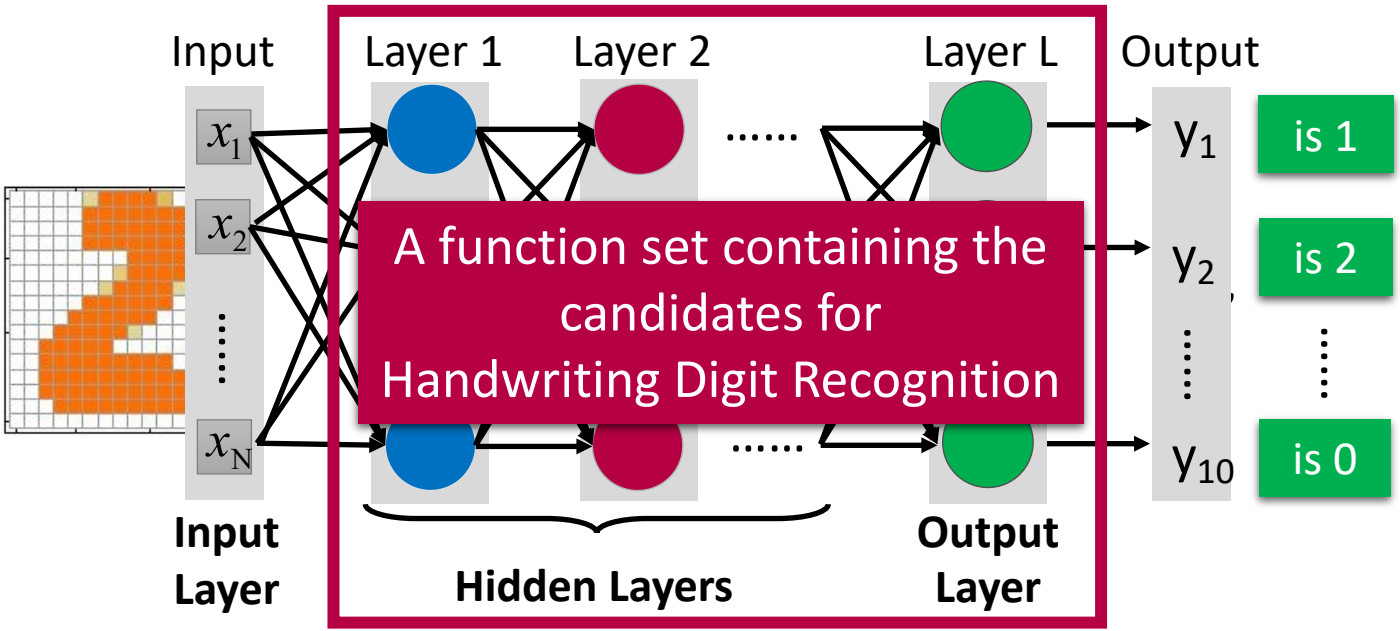

You need to decide the network structure to let a good function in your function set.

FAQ

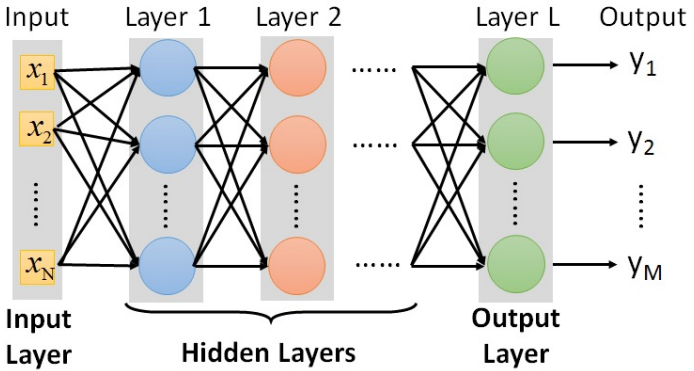

- Q: How many layers? How many neurons for each layer?

Trial and Error + Intuition

- Q: Can the structure be automatically determined?
  - E.g. Evolutionary Artificial Neural Networks
- Q: Can we design the network structure?

Convolutional Neural Network (CNN)

Training a Neural Network

Pick a network architecture (connectivity pattern between neurons)

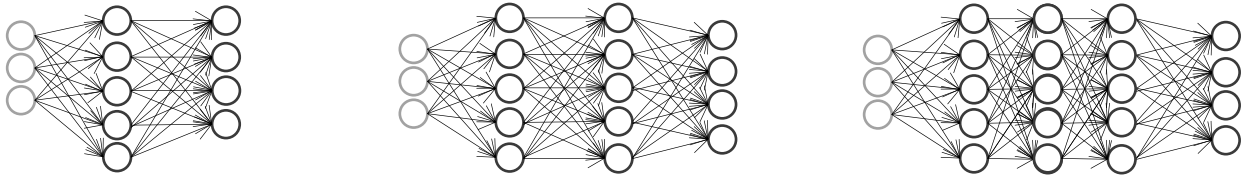

No. of input units: Dimension of features $x^{(i)}$

No. output units: Number of classes

Reasonable default: 1 hidden layer, or if >1 hidden layer, have same no. of hidden units in every layer (usually the more the better)

Three Steps for Deep Learning: Part II

Loss for an Example

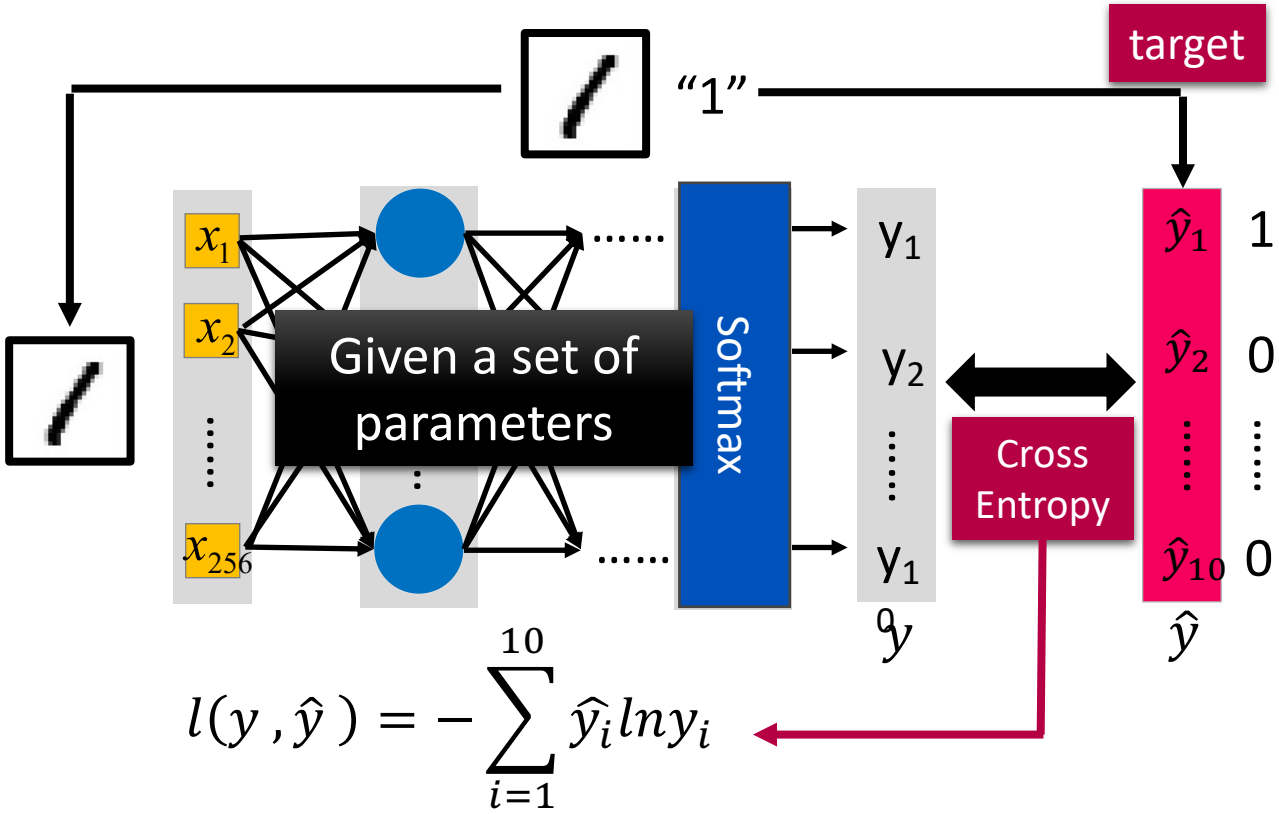

Total Loss

For all training data …

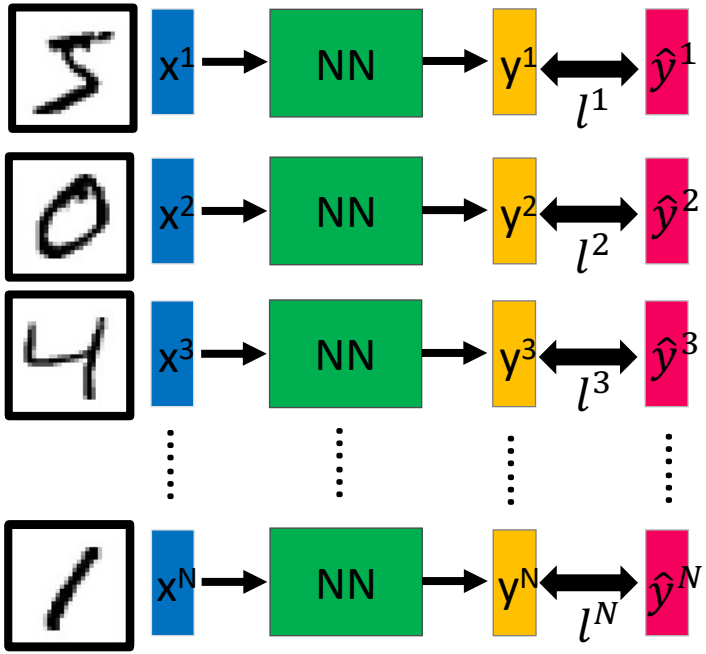

Total Loss:

$$L = \sum_{n=1}^{N} l^n$$

Find ***a function in function set*** that minimizes total loss L

Find ***the network parameters*** θ^* that minimize total loss L

Three Steps for Deep Learning: Part III

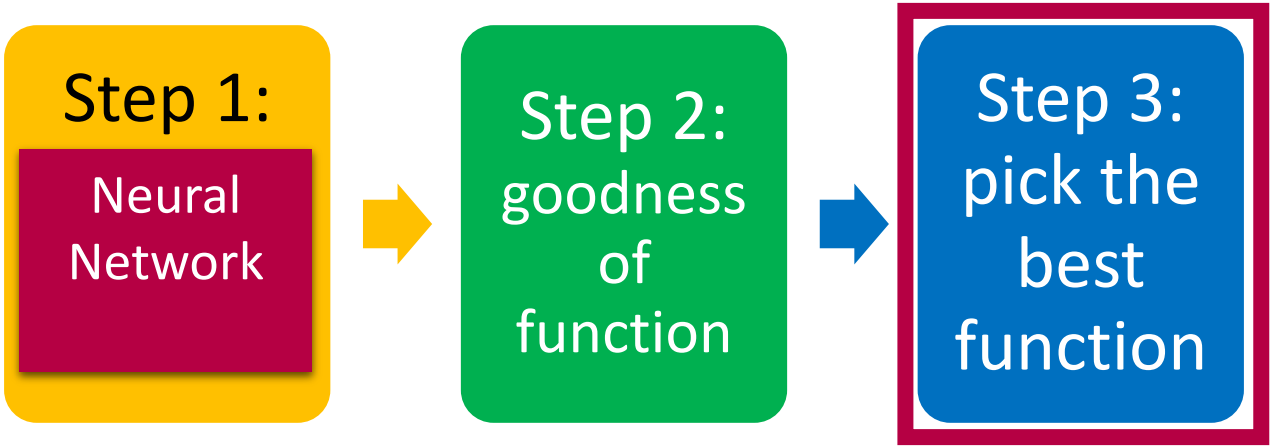

Gradient Descent

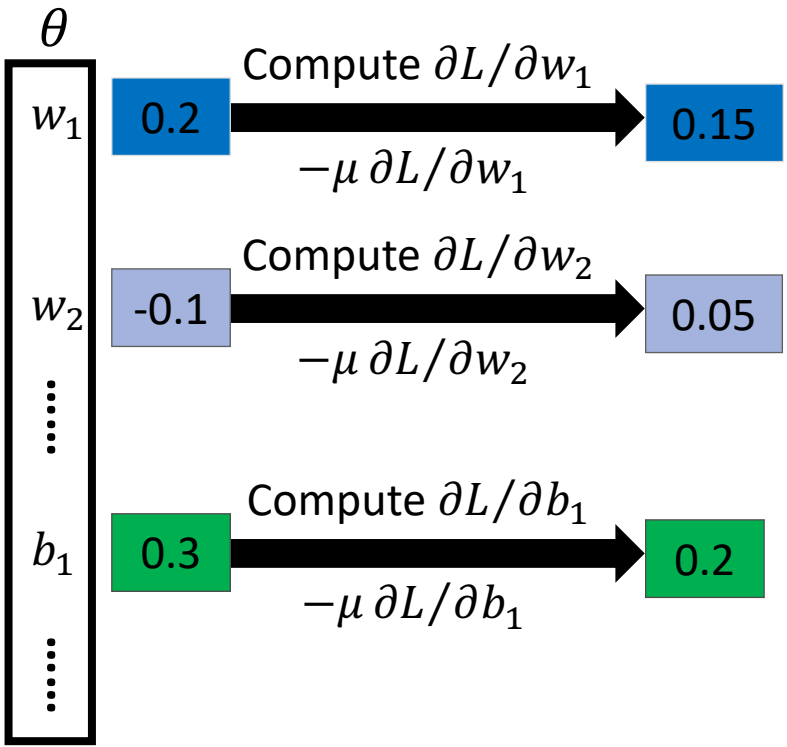

$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \\ \vdots \\ \frac{\partial L}{\partial b_1} \\ \vdots \end{bmatrix}$$

gradient

Gradient Descent, Continued

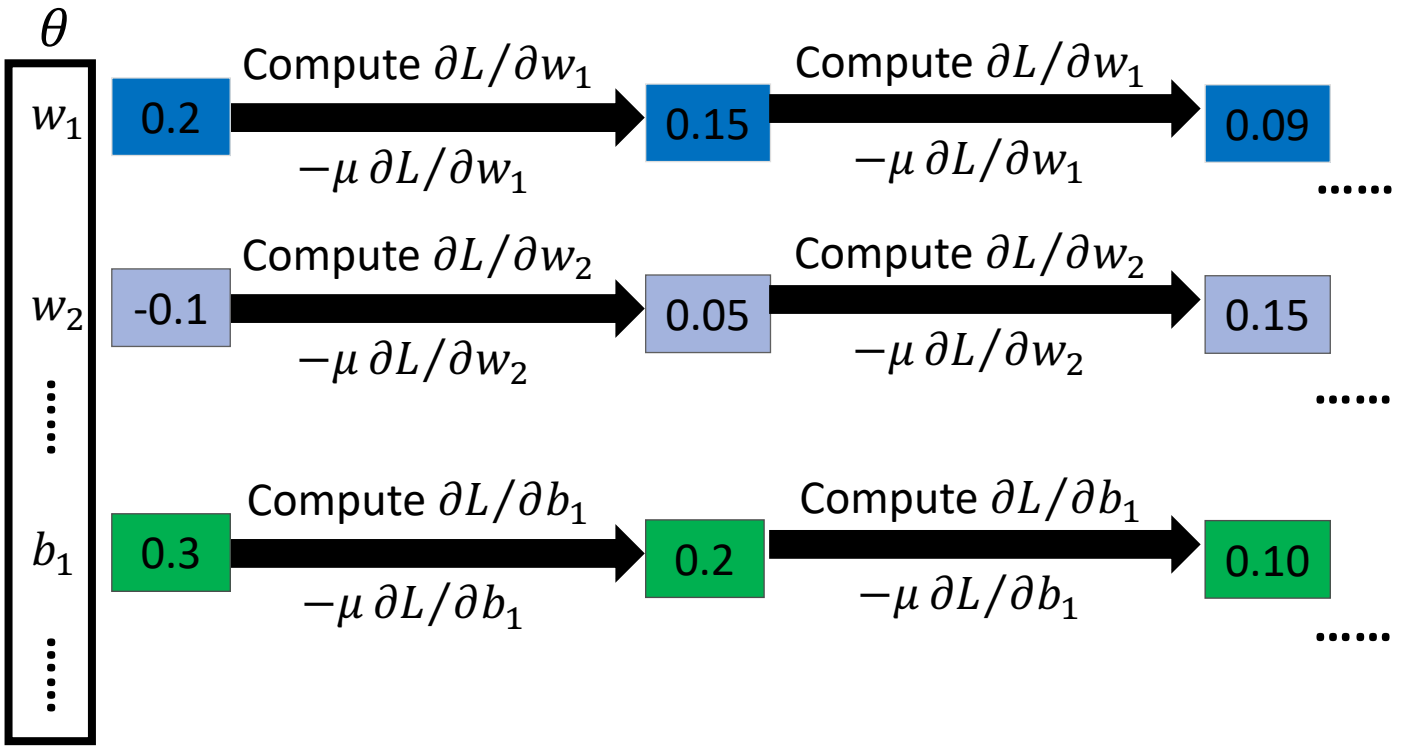

Backpropagation

Backpropagation: an efficient way to compute $\partial L/\partial w$ in neural network